

Harish G

AI Engineer | LLM, Agentic AI & Full-Stack

Bengaluru, India | +91 7892855850 | harishgreddy.work@gmail.com

LinkedIn | GitHub

Profile

AI Engineer with hands-on experience building production-grade LLM and agentic AI systems across voice and messaging channels. Skilled in prompt engineering, LLM evaluation, RAG and agent orchestration with tool calling and safe escalation, delivered through full-stack services (React, Node.js, Python / FastAPI). Built banking and healthcare assistants with role-based access control, audit logging and compliance-aware design, directly relevant to regulated environments. Strong foundation in data structures, system design and a production-engineering mindset: testing, CI/CD, observability and reliability.

Core Skills

AI / LLM: Prompt Engineering (system & tool prompts, guardrails, templates, versioning), LLM Evaluation (quality, accuracy, safety/bias, latency & cost, human-in-the-loop), Agentic AI (orchestration, tool calling, decision traces, safe escalation), RAG, Context Management, MCP (Model Context Protocol) concepts

Voice & Messaging: Real-time STT → LLM → TTS pipelines, telephony (Plivo, Twilio), WhatsApp, voice activity detection, multilingual handling

Languages: Python, JavaScript, TypeScript, Java, SQL, Kotlin

Frontend: React, Vite, Tailwind CSS

Backend: Node.js, Express, FastAPI, Spring Boot; REST API design, event-driven & WebSocket patterns

Data: PostgreSQL, MongoDB, Neo4j, Supabase / pgvector, Firebase

Cloud & DevOps: Docker, CI/CD, Git / GitLab, observability & monitoring, cloud-native deployment

Security & Controls: JWT authentication, role-based access control (RBAC), audit logging, data isolation, least-privilege design

Foundations: Data Structures & Algorithms, System Design, Distributed Systems concepts, modular / microservice design

Experience

Software Developer, CortexCraft.ai

Jan 2026 – Present

- Build production-grade AI features integrating LLMs, STT and TTS into real-world applications across voice and messaging channels.
- Design prompt frameworks (system & tool prompts, guardrails, templates) and apply RAG and context management for consistent, reliable agent behaviour.
- Implement agentic patterns: orchestration with tool calling, deterministic fallbacks and safe escalation paths.
- Iterate on AI quality, accuracy, latency and cost with an evaluation-driven approach, embedding engineering hygiene: testing, modular design and observability.

Freelance Web & App Developer (Self-Employed)

2023 – 2026

- Delivered full-stack web and app projects end-to-end for clients using React, Node.js and Python.
- Shipped responsive UIs, REST APIs and automation scripts with CI/CD and reliable deployments.

Selected Projects

FinGraph AI: Banking RAG Assistant (Regulated-Environment Design) *Node.js, Express, React, Neo4j, Supabase pgvector, JWT*

- Built a banking assistant over a Neo4j relationship graph with pgvector retrieval, plus query classification and workflow orchestration.
- Enforced security end-to-end: JWT auth (bcrypt), role-based access control (admin / manager / customer) applied before retrieval, and append-only audit logging of every query.
- Demonstrates compliance-aware patterns for regulated banking: least-privilege access, auditability and

safe failure modes.

Refyne Voice Agent: Real-Time Contact-Centre Voice AI *Python, FastAPI, Pipecat, Plivo, Sarvam, WebSocket, VAD*

- Built a real-time phone-call voice agent streaming STT → LLM → TTS over telephony, orchestrated with Pipecat.
- Low-latency turn-taking via voice activity detection, automatic language detection, and inbound + outbound call handling.
- Mirrors contact-centre concierge / triage and real-time assist patterns over voice channels.

KisanVoice AI: Multilingual Voice & Messaging Platform *Node.js, Express, React, MongoDB, Twilio / WhatsApp, Sarvam STT/TTS, Socket.io, Docker*

- WhatsApp-based voice survey platform with STT/TTS in 5 Indian languages and conditional conversation logic.
- Real-time admin dashboard over Socket.io with an audio QC workflow; containerised with Docker.

MediGraph AI: Hybrid RAG Clinical Assistant *Node.js, React, Neo4j, Supabase pgvector, HF Embeddings*

- Hybrid retrieval fusing a Neo4j knowledge graph with pgvector semantic search for grounded answers.
- Role-based access isolating patient and clinician data; interactive force-directed graph visualisation.

Zentrax: Agentic Desktop Assistant *Python, Whisper, Ollama, MediaPipe, WebSocket, Docker*

- LLM-based intent recognition (Ollama) with a modular tool / execution engine and fallback pipelines for reliability.
- Local LLM deployment via Docker for privacy-focused, on-device automation.

Highlights

- Built and deployed production-style AI and full-stack applications with real-user testing.
- Comfortable owning features end-to-end: design, implementation, observability and iteration.

Education

B.E., Information Science and Engineering

Vivekananda Institute of Technology, Bangalore